



Single insert – per-call overhead

Each impl inserts 300 docs one at a time; lower = faster




Bulk insert – throughput

Includes doc-object creation + batched insert; higher = faster



Search latency vs. collection size

Average over 100 queries (top-10); lower = faster



Overhead ratio vs. direct Python

1x = same speed as raw reticulate zvec_py; dashed line = baseline

